



Figure 1: Number of steps taken to find the goal in a randomly generated grid world maze of size  $50 \times 50$ . The light green area shows the range of the RBQL agent's highest and lowest step count, excluding the highest and lowest two, and the green line shows the average performance.