



Figure 1: Sampling algorithms choose the next token at each time step  $s_i$  by sampling from the conditional distribution  $p_\theta(\cdot | y_{:i-1})$  and appending it to the context.